

```
1 # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist\_cnn.py
2 from __future__ import print_function
3 import keras
4 from keras.datasets import mnist
5 from keras.models import Sequential
6 from keras.layers import Dense, Dropout, Flatten
7 from keras.layers import Conv2D, MaxPooling2D
8 from keras import backend as K
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 %matplotlib inline

1
2
3 batch_size = 128
4 num_classes = 10
5 epochs = 12
6
7 # input image dimensions
8 img_rows, img_cols = 28, 28
9
10 # the data, split between train and test sets
11 (x_train, y_train), (x_test, y_test) = mnist.load_data()
12
13 if K.image_data_format() == 'channels_first':
14     x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
15     x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
16     input_shape = (1, img_rows, img_cols)
17 else:
18     x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
19     x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
20     input_shape = (img_rows, img_cols, 1)
21
22 x_train = x_train.astype('float32')
23 x_test = x_test.astype('float32')
24 x_train /= 255
25 x_test /= 255
26 print('x_train shape:', x_train.shape)
27 print(x_train.shape[0], 'train samples')
28 print(x_test.shape[0], 'test samples')
29
30 # convert class vectors to binary class matrices
31 y_train = keras.utils.to_categorical(y_train, num_classes)
32 y_test = keras.utils.to_categorical(y_test, num_classes)
33
34 model = Sequential()
35 model.add(Conv2D(32, kernel_size=(3, 3),
36                 activation='relu',
37                 input_shape=input_shape))
38 model.add(Conv2D(64, (3, 3), activation='relu'))
39 model.add(MaxPooling2D(pool_size=(2, 2)))
40 model.add(Dropout(0.25))
41 model.add(Flatten())
```

```
42 model.add(Dense(128, activation='relu'))
43 model.add(Dropout(0.5))
44 model.add(Dense(num_classes, activation='softmax'))
45
46 model.compile(loss=keras.losses.categorical_crossentropy,
47               optimizer=keras.optimizers.Adadelta(),
48               metrics=['accuracy'])
49
50 model.fit(x_train, y_train,
51         batch_size=batch_size,
52         epochs=epochs,
53         verbose=1,
54         validation_data=(x_test, y_test))
55 score = model.evaluate(x_test, y_test, verbose=0)
56 print('Test loss:', score[0])
57 print('Test accuracy:', score[1])
```



```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - ke
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizer
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
Train on 60000 samples, validate on 10000 samples
Epoch 1/12
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t
60000/60000 [=====] - 20s 338us/step - loss: 0.2541 -
Epoch 2/12
60000/60000 [=====] - 6s 100us/step - loss: 0.0840 -
Epoch 3/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0637 -
Epoch 4/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0518 -
Epoch 5/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0443 -
Epoch 6/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0405 -
Epoch 7/12
60000/60000 [=====] - 6s 104us/step - loss: 0.0355 -
Epoch 8/12
60000/60000 [=====] - 6s 101us/step - loss: 0.0327 -
Epoch 9/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0296 -
Epoch 10/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0262 -
Epoch 11/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0254 -
Epoch 12/12
60000/60000 [=====] - 6s 102us/step - loss: 0.0256 -
Test loss: 0.027925510946988835
Test accuracy: 0.9905

```

▼ ===== Assignment =====

```

1 import numpy as np
2 import time
3 # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
4 # https://stackoverflow.com/a/14434334
5 # this function is used to update the plots for each epoch and error
6 def plt_dynamic(x, vy, ty, ax, colors=['b']):
7     ax.plot(x, vy, 'b', label="Validation Loss")
8     ax.plot(x, ty, 'r', label="Train Loss")
9     plt.legend()
10    plt.grid()
11    fig.canvas.draw()

```

Model 1: 7 conv layers + filter size = 3x3 + relu activation + d

```

1 model = Sequential()
2
3 model.add(Conv2D(20, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
4 model.add(MaxPooling2D(pool_size=(2, 2)))
5 model.add(Dropout(0.25))
6
7 model.add(Conv2D(25, (3, 3), activation='relu', padding='same'))
8 model.add(MaxPooling2D(pool_size=(2, 2)))
9 model.add(Dropout(0.25))
10
11 model.add(Conv2D(35, (3, 3), activation='relu', padding='same'))
12 model.add(MaxPooling2D(pool_size=(2, 2)))
13
14 model.add(Conv2D(45, (3, 3), activation='relu', padding='same'))
15 model.add(MaxPooling2D(pool_size=(2, 2)))
16
17 model.add(Conv2D(55, (3, 3), activation='relu', padding='same'))
18
19
20 model.add(Conv2D(65, (3, 3), activation='relu', padding='same'))
21

```

```

22
23 model.add(Conv2D(75, (3, 3), activation='relu',padding='same'))
24 model.add(Dropout(0.75))
25
26 model.add(Flatten())
27 model.add(Dense(128, activation='relu'))
28 model.add(Dropout(0.5))
29 model.add(Dense(num_classes, activation='softmax'))
30
31 model.compile(loss=keras.losses.categorical_crossentropy,
32               optimizer=keras.optimizers.Adadelta(),
33               metrics=['accuracy'])
34
35 history = model.fit(x_train, y_train,
36                    batch_size=batch_size,
37                    epochs=epochs,
38                    verbose=1,
39                    validation_data=(x_test, y_test))
40 score = model.evaluate(x_test, y_test, verbose=0)
41 print('Test loss:', score[0])
42 print('Test accuracy:', score[1])
43

```

⏏ WARNING:tensorflow:Large dropout rate: 0.75 (>0.5). In TensorFlow 2.x, dropout Train on 60000 samples, validate on 10000 samples

```

Epoch 1/12
60000/60000 [=====] - 6s 92us/step - loss: 1.9629 - a
Epoch 2/12
60000/60000 [=====] - 5s 77us/step - loss: 0.8247 - a
Epoch 3/12
60000/60000 [=====] - 5s 75us/step - loss: 0.4365 - a
Epoch 4/12
60000/60000 [=====] - 5s 75us/step - loss: 0.3244 - a
Epoch 5/12
60000/60000 [=====] - 5s 76us/step - loss: 0.2699 - a
Epoch 6/12
60000/60000 [=====] - 4s 74us/step - loss: 0.2348 - a
Epoch 7/12
60000/60000 [=====] - 4s 74us/step - loss: 0.2188 - a
Epoch 8/12
60000/60000 [=====] - 5s 75us/step - loss: 0.1909 - a
Epoch 9/12
60000/60000 [=====] - 5s 77us/step - loss: 0.1835 - a
Epoch 10/12
60000/60000 [=====] - 5s 75us/step - loss: 0.1739 - a
Epoch 11/12
60000/60000 [=====] - 5s 76us/step - loss: 0.1600 - a
Epoch 12/12
60000/60000 [=====] - 5s 77us/step - loss: 0.1572 - a
Test loss: 0.07009304843556602
Test accuracy: 0.986

```

```

1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig, ax = plt.subplots(1, 1)


```

```

5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,12+1))
10 # we will get val_loss and val_acc only when you pass the paramter validation_d
11 # val_loss : validation loss
12 # val_acc : validation accuracy
13
14 # loss : training loss
15 # acc : train accuracy
16 # for each key in history.history we will have a list of length equal to numb
17
18 vy = history.history['val_loss']
19 ty = history.history['loss']
20 plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07009304843556602
Test accuracy: 0.986



```


1 w_after = model.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='v')

```

```

22 ax = sns.violinplot(y=out_w,color=y )
23 plt.xlabel('Output Layer ')
24 plt.show()

```



Model 2: 5 conv layers + filter size = 3x3 + relu activation + d

```

1 model = Sequential()
2
3 model.add(Conv2D(40, kernel_size=(3, 3),activation='relu',input_shape=input_sha
4 model.add(MaxPooling2D(pool_size=(2, 2)))
5 model.add(Dropout(0.25))
6
7 model.add(Conv2D(40, (3, 3), activation='relu',padding='same'))
8 model.add(MaxPooling2D(pool_size=(2, 2)))
9 model.add(Dropout(0.25))
10
11 model.add(Conv2D(50, (3, 3), activation='relu',padding='same'))
12 model.add(MaxPooling2D(pool_size=(2, 2)))
13
14 model.add(Conv2D(55, (3, 3), activation='relu',padding='same'))
15 #model.add(MaxPooling2D(pool_size=(2, 2)))
16
17 model.add(Conv2D(55, (3, 3), activation='relu',padding='same'))
18 #model.add(MaxPooling2D(pool_size=(2, 2)))
19
20
21 model.add(Flatten())
22 model.add(Dense(128, activation='relu'))

```

```

22 model.add(Dense(128, activation='relu'))
23 model.add(Dropout(0.5))
24 model.add(Dense(num_classes, activation='softmax'))
25
26 model.summary()

```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
conv2d_37 (Conv2D)	(None, 26, 26, 40)	400
max_pooling2d_27 (MaxPooling)	(None, 13, 13, 40)	0
dropout_21 (Dropout)	(None, 13, 13, 40)	0
conv2d_38 (Conv2D)	(None, 13, 13, 40)	14440
max_pooling2d_28 (MaxPooling)	(None, 6, 6, 40)	0
dropout_22 (Dropout)	(None, 6, 6, 40)	0
conv2d_39 (Conv2D)	(None, 6, 6, 50)	18050
max_pooling2d_29 (MaxPooling)	(None, 3, 3, 50)	0
conv2d_40 (Conv2D)	(None, 3, 3, 55)	24805
conv2d_41 (Conv2D)	(None, 3, 3, 55)	27280
flatten_6 (Flatten)	(None, 495)	0
dense_11 (Dense)	(None, 128)	63488
dropout_23 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 10)	1290
Total params: 149,753		
Trainable params: 149,753		
Non-trainable params: 0		

```

1 model.compile(loss=keras.losses.categorical_crossentropy,
2               optimizer=keras.optimizers.Adadelta(),
3               metrics=['accuracy'])
4
5 history = model.fit(x_train, y_train,
6                   batch_size=batch_size,
7                   epochs=epochs,
8                   verbose=1

```



```

8         verbose=1,
9         validation_data=(x_test, y_test))
10 score = model.evaluate(x_test, y_test, verbose=0)
11 print('Test loss:', score[0])
12 print('Test accuracy:', score[1])

```

☞ Train on 60000 samples, validate on 10000 samples

```

Epoch 1/12
60000/60000 [=====] - 6s 104us/step - loss: 0.4544 - a
Epoch 2/12
60000/60000 [=====] - 5s 83us/step - loss: 0.1046 - a
Epoch 3/12
60000/60000 [=====] - 5s 83us/step - loss: 0.0747 - a
Epoch 4/12
60000/60000 [=====] - 5s 85us/step - loss: 0.0580 - a
Epoch 5/12
60000/60000 [=====] - 5s 83us/step - loss: 0.0510 - a
Epoch 6/12
60000/60000 [=====] - 5s 83us/step - loss: 0.0437 - a
Epoch 7/12
60000/60000 [=====] - 5s 83us/step - loss: 0.0374 - a
Epoch 8/12
60000/60000 [=====] - 5s 86us/step - loss: 0.0343 - a
Epoch 9/12
60000/60000 [=====] - 5s 82us/step - loss: 0.0311 - a
Epoch 10/12
60000/60000 [=====] - 5s 85us/step - loss: 0.0293 - a
Epoch 11/12
60000/60000 [=====] - 5s 82us/step - loss: 0.0269 - a
Epoch 12/12
60000/60000 [=====] - 5s 83us/step - loss: 0.0255 - a
Test loss: 0.024910171242001615
Test accuracy: 0.9927

```

```

1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,12+1))
10 # we will get val_loss and val_acc only when you pass the paramter validation_d
11 # val_loss : validation loss
12 # val_acc : validation accuracy
13
14 # loss : training loss
15 # acc : train accuracy
16 # for each key in history history we will have a list of length equal to numh


```

```

16 # For each key in history.history we will have a list of length equal to number of epochs
17
18 vy = history.history['val_loss']
19 ty = history.history['loss']
20 plt_dynamic(x, vy, ty, ax)

```


Test score: 0.024910171242001615
Test accuracy: 0.9927



```

1 w_after = model.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```



Model 3: 3 Convolution Layers + filter size = 5x5 + dropout +

```

1 model = Sequential()
2
3 model.add(Conv2D(40, kernel_size=(5, 5),activation='relu',input_shape=input_sha
4 model.add(MaxPooling2D(pool_size=(2, 2)))
5 model.add(Dropout(0.25))
6
7 model.add(Conv2D(50, (5, 5), activation='relu',padding='same'))
8
9
10 model.add(Conv2D(55, (5, 5), activation='relu',padding='same'))
11
12 model.add(Flatten())
13 model.add(Dense(128, activation='relu'))
14 model.add(Dropout(0.5))
15 model.add(Dense(num_classes, activation='softmax'))
16
17 model.summary()

```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d_42 (Conv2D)	(None, 24, 24, 40)	1040
max_pooling2d_30 (MaxPooling)	(None, 12, 12, 40)	0
dropout_24 (Dropout)	(None, 12, 12, 40)	0
conv2d_43 (Conv2D)	(None, 12, 12, 50)	50050
conv2d_44 (Conv2D)	(None, 12, 12, 55)	68805
flatten_7 (Flatten)	(None, 7920)	0
dense_13 (Dense)	(None, 128)	1013888
dropout_25 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 10)	1290
Total params: 1,135,073		
Trainable params: 1,135,073		
Non-trainable params: 0		

```

1 model.compile(loss=keras.losses.categorical_crossentropy,
2               optimizer=keras.optimizers.Adadelta(),
3               metrics=['accuracy'])
4
5 history = model.fit(x_train, y_train,
6                   batch_size=batch_size,
7                   epochs=epochs,
8                   verbose=1,
9                   validation_data=(x_test, y_test))
10 score = model.evaluate(x_test, y_test, verbose=0)
11 print('Test loss:', score[0])
12 print('Test accuracy:', score[1])

```



```

Train on 60000 samples, validate on 10000 samples
Epoch 1/12
60000/60000 [=====] - 7s 118us/step - loss: 0.2478 - a
Epoch 2/12
60000/60000 [=====] - 6s 99us/step - loss: 0.0675 - a
Epoch 3/12
60000/60000 [=====] - 6s 101us/step - loss: 0.0445 - a
Epoch 4/12
60000/60000 [=====] - 6s 99us/step - loss: 0.0356 - a
Epoch 5/12
60000/60000 [=====] - 6s 100us/step - loss: 0.0286 - a
Epoch 6/12
60000/60000 [=====] - 6s 101us/step - loss: 0.0228 - a
Epoch 7/12
60000/60000 [=====] - 6s 99us/step - loss: 0.0205 - a
Epoch 8/12
60000/60000 [=====] - 6s 105us/step - loss: 0.0178 - a
Epoch 9/12
60000/60000 [=====] - 6s 103us/step - loss: 0.0159 - a
Epoch 10/12
60000/60000 [=====] - 6s 101us/step - loss: 0.0137 - a
Epoch 11/12
60000/60000 [=====] - 6s 99us/step - loss: 0.0125 - a
Epoch 12/12
60000/60000 [=====] - 6s 99us/step - loss: 0.0117 - a
Test loss: 0.022570345474630403
Test accuracy: 0.9934

```

```

1 score = model.evaluate(x_test, y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,12+1))
10 # we will get val_loss and val_acc only when you pass the paramter validation_d
11 # val_loss : validation loss
12 # val_acc : validation accuracy
13
14 # loss : training loss
15 # acc : train accuracy
16 # for each key in history history we will have a list of length equal to numbr


```

```

16 # For each key in history.history we will have a list of length equal to numba
17
18 vy = history.history['val_loss']
19 ty = history.history['loss']
20 plt_dynamic(x, vy, ty, ax)

```


Test score: 0.022570345474630403
Test accuracy: 0.9934



```

1 w_after = model.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```



Conclusion

```
1 from prettytable import PrettyTable
2 x = PrettyTable()
3 x.field_names = ["Model", "Conv Layers", "Filter Size", "Max Pool", "Test Accuracy"]
4 x.add_row(["1", "7", "(3X3)", "(2x2)", "98.60"])
5 x.add_row(["2", "5", "(3x3)", "(2x2)", "99.27"])
6 x.add_row(["3", "3", "(5x5)", "(2x2)", "99.34"])
7 print(x)
```

Model	Conv Layers	Filter Size	Max Pool	Test Accuracy
1	7	(3X3)	(2x2)	98.60
2	5	(3x3)	(2x2)	99.27
3	3	(5x5)	(2x2)	99.34

