

▼ Keras -- MLPs on MNIST

```
1 # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow
2 from keras.utils import np_utils
3 from keras.datasets import mnist
4 import seaborn as sns
5 from keras.initializers import RandomNormal
6 from keras.initializers import he_normal
7 from keras.layers.normalization import BatchNormalization
8 from keras.layers import Dropout
```

```
1 %matplotlib notebook
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import time
5 # https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
6 # https://stackoverflow.com/a/14434334
7 # this function is used to update the plots for each epoch and error
8 def plt_dynamic(x, vy, ty, ax, colors=['b']):
9     ax.plot(x, vy, 'b', label="Validation Loss")
10    ax.plot(x, ty, 'r', label="Train Loss")
11    plt.legend()
12    plt.grid()
13    fig.canvas.draw()
```

```
1 # the data, shuffled and split between train and test sets
2 (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

📄 Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```
1 print("Number of training examples :", X_train.shape[0], "and each image is of
2 print("Number of test examples :", X_test.shape[0], "and each image is of shape
```

📄 Number of training examples : 60000 and each image is of shape (28, 28)
Number of test examples : 10000 and each image is of shape (28, 28)

```
1 # if you observe the input shape its 2 dimensional vector
2 # for each image we have a (28*28) vector
3 # we will convert the (28*28) vector into single dimensional vector of 1 * 784
4
5 X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
6 X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
1 # after converting the input images from 3d to 2d vectors
2
3 print("Number of training examples :", X_train.shape[0], "and each image is of
4 print("Number of training examples :", X_test.shape[0], "and each image is of s
```

```

☞ Number of training examples : 60000 and each image is of shape (784)
   Number of training examples : 10000 and each image is of shape (784)

```

```

1 # An example data point
2 print(X_train[0])

```

```

☞ [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  3  18  18  18 126 136 175 26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30 36 94 154
170 253 253 253 253 253 225 172 253 242 195 64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93 82
 82 56 39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 80 156 107 253 253 205 11  0 43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190 2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0 81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 249 253 249 64  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0 39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0 24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]

```

```


1 # if we observe the above matrix each cell is having a value between 0-255
2 # before we move to apply machine learning algorithms lets try to normalize the
3 # X => (X - Xmin)/(Xmax-Xmin) = X/255
4
5 X_train = X_train/255
6 X_test = X_test/255

```

```
1 # example data point after normlizing  
2 print(X_train[0])
```



1



```

1 # here we are having a class number for each image
2 print("Class label of first image :", y_train[0])
3
4 # lets convert this into a 10 dimensional vector
5 # ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
6 # this conversion needed for MLPs
7
8 Y_train = np_utils.to_categorical(y_train, 10)
9 Y_test = np_utils.to_categorical(y_test, 10)
10
11 print("After converting the output into a vector : ",Y_train[0])

```

 Class label of first image : 5
 After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

Softmax classifier

```

1 # https://keras.io/getting-started/sequential-model-guide/
2
3 # The Sequential model is a linear stack of layers.
4 # you can create a Sequential model by passing a list of layer instances to the
5
6 # model = Sequential([
7 #     Dense(32, input_shape=(784,)),
8 #     Activation('relu'),
9 #     Dense(10),
10 #     Activation('softmax'),
11 # ])
12
13 # You can also simply add layers via the .add() method:
14
15 # model = Sequential()
16 # model.add(Dense(32, input_dim=784))
17 # model.add(Activation('relu'))
18
19 ###
20
21 # https://keras.io/layers/core/
22
23 # keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer=
24 # bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, act
25 # kernel_constraint=None, bias_constraint=None)
26
27 # Dense implements the operation: output = activation(dot(input, kernel) + bias

```

```

28 # activation is the element-wise activation function passed as the activation a
29 # kernel is a weights matrix created by the layer, and
30 # bias is a bias vector created by the layer (only applicable if use_bias is Tr
31
32 # output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)
33
34 #####
35
36 # https://keras.io/activations/
37
38 # Activations can either be used through an Activation layer, or through the ac
39
40 # from keras.layers import Activation, Dense
41
42 # model.add(Dense(64))
43 # model.add(Activation('tanh'))
44
45 # This is equivalent to:
46 # model.add(Dense(64, activation='tanh'))
47
48 # there are many activation functions ar available ex: tanh, relu, softmax
49
50
51 from keras.models import Sequential
52 from keras.layers import Dense, Activation
53
54
55 1 # some model parameters
56 2
57 3 output_dim = 10
58 4 input_dim = X_train.shape[1]
59 5
60 6 batch_size = 128
61 7 nb_epoch = 20
62
63
64 1 # start building a model
65 2 model = Sequential()
66 3
67 4 # The model needs to know what input shape it should expect.
68 5 # For this reason, the first layer in a Sequential model
69 6 # (and only the first, because following layers can do automatic shape inferenc
70 7 # needs to receive information about its input shape.
71 8 # you can use input_shape and input_dim to pass the shape of input
72 9
73 10 # output_dim represent the number of nodes need in that layer
74 11 # here we have 10 nodes
75 12
76 13 model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tf

```
1 # Before training a model, you need to configure the learning process, which is
2
3 # It receives three arguments:
4 # An optimizer. This could be the string identifier of an existing optimizer ,
5 # A loss function. This is the objective that the model will try to minimize.,
6 # A list of metrics. For any classification problem you will want to set this to
7
8
9 # Note: when using the categorical_crossentropy loss, your targets should be in
10 # (e.g. if you have 10 classes, the target for each sample should be a 10-dimen
11 # for a 1 at the index corresponding to the class of the sample).
12
13 # that is why we converted out labels into vectors
14
15 model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
16
17 # Keras models are trained on Numpy arrays of input data and labels.
18 # For training a model, you will typically use the fit function
19
20 # fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
21 # validation_data=None, shuffle=True, class_weight=None, sample_weight=None, in
22 # validation_steps=None)
23
24 # fit() function Trains the model for a fixed number of epochs (iterations on a
25
26 # it returns A History object. Its History.history attribute is a record of tra
27 # metrics values at successive epochs, as well as validation loss values and va
28
29 # https://github.com/openai/baselines/issues/20
30
31 history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, va
32
```




```

Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/t


60000/60000 [=====] - 11s 176us/step - loss: 1.2651 -
Epoch 2/20
60000/60000 [=====] - 1s 24us/step - loss: 0.7152 - a
Epoch 3/20
60000/60000 [=====] - 1s 24us/step - loss: 0.5881 - a
Epoch 4/20
60000/60000 [=====] - 1s 25us/step - loss: 0.5267 - a
Epoch 5/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4892 - a
Epoch 6/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4633 - a
Epoch 7/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4441 - a
Epoch 8/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4291 - a
Epoch 9/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4170 - a
Epoch 10/20
60000/60000 [=====] - 1s 24us/step - loss: 0.4069 - a
Epoch 11/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3984 - a
Epoch 12/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3910 - a
Epoch 13/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3846 - a
Epoch 14/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3789 - a
Epoch 15/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3738 - a
Epoch 16/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3693 - a
Epoch 17/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3651 - a
Epoch 18/20
60000/60000 [=====] - 2s 25us/step - loss: 0.3614 - a
Epoch 19/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3579 - a
Epoch 20/20
60000/60000 [=====] - 1s 23us/step - loss: 0.3547 - a

```

```

1 score = model.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_d
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to numb
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

 Test score: 0.33514436384439467
 Test accuracy: 0.909

MLP + Sigmoid activation + SGDOptimizer

```

1 # Multilayer perceptron
2
3 model_sigmoid = Sequential()
4 model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
5 model_sigmoid.add(Dense(128, activation='sigmoid'))
6 model_sigmoid.add(Dense(output_dim, activation='softmax'))
7
8 model_sigmoid.summary()

```



```
Model: "sequential_2"
```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920

```
1 model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=
```

```
2
```

```
3 history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
```

```
☞ Train on 60000 samples, validate on 10000 samples
```

```
Epoch 1/20
```

```
60000/60000 [=====] - 2s 30us/step - loss: 2.2663 - a
```

```
Epoch 2/20
```

```
60000/60000 [=====] - 2s 26us/step - loss: 2.1799 - a
```

```
Epoch 3/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 2.0670 - a
```

```
Epoch 4/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 1.9053 - a
```

```
Epoch 5/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 1.6921 - a
```

```
Epoch 6/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 1.4557 - a
```

```
Epoch 7/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 1.2397 - a
```

```
Epoch 8/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 1.0686 - a
```

```
Epoch 9/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 0.9402 - a
```

```
Epoch 10/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 0.8436 - a
```

```
Epoch 11/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 0.7690 - a
```

```
Epoch 12/20
```

```
60000/60000 [=====] - 2s 26us/step - loss: 0.7099 - a
```

```
Epoch 13/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 0.6624 - a
```

```
Epoch 14/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 0.6232 - a
```

```
Epoch 15/20
```

```
60000/60000 [=====] - 2s 26us/step - loss: 0.5907 - a
```

```
Epoch 16/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 0.5633 - a
```

```
Epoch 17/20
```

```
60000/60000 [=====] - 2s 27us/step - loss: 0.5398 - a
```

```
Epoch 18/20
```

```
60000/60000 [=====] - 2s 26us/step - loss: 0.5195 - a
```

```
Epoch 19/20
```

```
60000/60000 [=====] - 2s 26us/step - loss: 0.5018 - a
```

```
Epoch 20/20
```

```
60000/60000 [=====] - 2s 26us/step - loss: 0.4862 - a
```

```
1 score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
```

```
2 print('Test score:', score[0])
```

```
3 print('Test accuracy:', score[1])
```

```
4
```

```
5 fig,ax = plt.subplots(1,1)
```

```
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
7
```

```
8 # List of epoch numbers
```

```

8 # list of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_d
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to numb
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```

☞ Test score: 0.4634165374755859
Test accuracy: 0.8754

```

1 w_after = model_sigmoid.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```

☞

MLP + Sigmoid activation + ADAM

```

1 model_sigmoid = Sequential()
2 model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))

```

```
3 model_sigmoid.add(Dense(128, activation='sigmoid'))
4 model_sigmoid.add(Dense(output_dim, activation='softmax'))
5
6 model_sigmoid.summary()
7
8 model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metric
9
10 history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_
```



Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920

```

1 score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_d
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to numb
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```



Test score: 0.06385514608082886

Test accuracy: 0.9824

```

1 w_after = model_sigmoid.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```



/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarning
 kde_data = remove_na(group_data)
 /usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarning
 violin_data = remove_na(group_data)

MLP + ReLU +SGD

```

1 # Multilayer perceptron
2
3 # https://arxiv.org/pdf/1707.09725.pdf#page=95
4 # for relu layers
5 # If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition:
6 #  $h1 \Rightarrow \sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0, \sigma) = N(0, 0.062)$ 
7 #  $h2 \Rightarrow \sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0, \sigma) = N(0, 0.125)$ 
8 #  $out \Rightarrow \sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 
9
10 model_relu = Sequential()
11 model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0, std=0.062)))
12 model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0, std=0.125)))
13 model_relu.add(Dense(output_dim, activation='softmax'))
14
15 model_relu.summary()

```



Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
1 model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
2
3 history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)
```




```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 67us/step - loss: 0.7579 - a
Epoch 2/20
60000/60000 [=====] - 4s 64us/step - loss: 0.3535 - a
Epoch 3/20
60000/60000 [=====] - 4s 64us/step - loss: 0.2900 - a

1 score = model_relu.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_d
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.histrory we will have a list of length equal to numb
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```



Test score: 0.12405014228336513

Test accuracy: 0.9631



```
1 w_after = model_relu.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()
```



Trained model Weights Trained model Weights Trained model Weights

MLP + ReLU + ADAM

```

1 model_relu = Sequential()
2 model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_i
3 model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(me
4 model_relu.add(Dense(output_dim, activation='softmax'))
5
6 print(model_relu.summary())
7
8 model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[
9
10 history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epo

```



Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290

Total params: 468,874

Trainable params: 468,874

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 121us/step - loss: 0.2341 - a

Epoch 2/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0878 - a

Epoch 3/20

60000/60000 [=====] - 5s 75us/step - loss: 0.0544 - a

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0354 - a

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0266 - a

Epoch 6/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0200 - a

Epoch 7/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0155 - a

Epoch 8/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0140 - a

Epoch 9/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0143 - a

Epoch 10/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0128 - a

Epoch 11/20

60000/60000 [=====] - 7s 125us/step - loss: 0.0081 - a

Epoch 12/20

60000/60000 [=====] - 8s 129us/step - loss: 0.0121 - a

Epoch 13/20

60000/60000 [=====] - 8s 133us/step - loss: 0.0107 - a

Epoch 14/20

60000/60000 [=====] - 8s 129us/step - loss: 0.0113 - a

Epoch 15/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0058 - a

Epoch 16/20

60000/60000 [=====] - 4s 65us/step - loss: 0.0040 - a

Epoch 17/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0119 - a

Epoch 18/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0105 - a

Epoch 19/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0064 - a

Epoch 20/20

60000/60000 [=====] - 4s 72us/step - loss: 0.0056 - a

```
1 score = model_relu.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_d
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in history.history we will have a list of length equal to numb
22
23
24 vy = history.history['val_loss']
25 ty = history.history['loss']
26 plt_dynamic(x, vy, ty, ax)
```




Test score: 0.10294274219236926

Test accuracy: 0.9805

```

1 w_after = model_relu.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```



MLP + Batch-Norm on hidden Layers + AdamOptimizer

```

1 # Multilayer perceptron
2
3 # https://intoli.com/blog/neural-network-initialization/
4 # If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition:
5 #  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
6 #  $h2 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
7 #  $h1 \Rightarrow \sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 
8
9 from keras.layers.normalization import BatchNormalization
10
11 model_batch = Sequential()
12
13 model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(0, 0.039)))
14 model_batch.add(BatchNormalization())
15
16 model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(0, 0.055)))
17 model_batch.add(BatchNormalization())
18
19 model_batch.add(Dense(output_dim, activation='softmax'))
20
21
22 model_batch.summary()

```



Layer (type)	Output Shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

```

1 model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
2
3 history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs)

```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 138us/step - loss: 0.3036 -
Epoch 2/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1747 -
Epoch 3/20
60000/60000 [=====] - 13s 220us/step - loss: 0.1367 -
Epoch 4/20
60000/60000 [=====] - 9s 156us/step - loss: 0.1134 -
Epoch 5/20
60000/60000 [=====] - 13s 211us/step - loss: 0.0949 -
Epoch 6/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0802 -
Epoch 7/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0682 -
Epoch 8/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0608 -
Epoch 9/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0532 -
Epoch 10/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0455 -
Epoch 11/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0376 -
Epoch 12/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0350 -
Epoch 13/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0308 -
Epoch 14/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0271 -
Epoch 15/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0231 -
Epoch 16/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0220 -
Epoch 17/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0229 -
Epoch 18/20
60000/60000 [=====] - 8s 128us/step - loss: 0.0203 -
Epoch 19/20
60000/60000 [=====] - 7s 125us/step - loss: 0.0171 -
Epoch 20/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0146 -
```

```
1 score = model_batch.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
```




```

11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_d
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in histry.history we will have a list of length equal to numb
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```



Test score: 0.10456635547156475
Test accuracy: 0.9732



```


1 w_after = model_batch.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)

```

```

11 plt.title("Trained model weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```



5. MLP + Dropout + AdamOptimizer

```

1 # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-layer
2
3 from keras.layers import Dropout
4
5 model_drop = Sequential()
6
7 model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer='he_normal'))
8 model_drop.add(BatchNormalization())
9 model_drop.add(Dropout(0.5))
10

```

```

11 model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal
12 model_drop.add(BatchNormalization())
13 model_drop.add(Dropout(0.5))
14
15 model_drop.add(Dense(output_dim, activation='softmax'))
16
17
18 model_drop.summary()

```



Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290

Total params: 471,434
 Trainable params: 470,154
 Non-trainable params: 1,280

```

1 model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[
2
3 history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epo

```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 14s 227us/step - loss: 0.6612 -
Epoch 2/20
60000/60000 [=====] - 8s 136us/step - loss: 0.4250 -
Epoch 3/20
60000/60000 [=====] - 12s 198us/step - loss: 0.3841 -
Epoch 4/20
60000/60000 [=====] - 8s 138us/step - loss: 0.3551 -
Epoch 5/20
60000/60000 [=====] - 7s 123us/step - loss: 0.3355 -
Epoch 6/20
60000/60000 [=====] - 8s 136us/step - loss: 0.3234 -
Epoch 7/20
60000/60000 [=====] - 8s 131us/step - loss: 0.3068 -
Epoch 8/20
60000/60000 [=====] - 11s 185us/step - loss: 0.2933 -
Epoch 9/20
60000/60000 [=====] - 13s 222us/step - loss: 0.2850 -
Epoch 10/20
60000/60000 [=====] - 14s 236us/step - loss: 0.2715 -
Epoch 11/20
60000/60000 [=====] - 8s 141us/step - loss: 0.2611 -
Epoch 12/20
60000/60000 [=====] - 8s 134us/step - loss: 0.2464 -
Epoch 13/20
60000/60000 [=====] - 8s 137us/step - loss: 0.2382 -
Epoch 14/20
60000/60000 [=====] - 8s 136us/step - loss: 0.2275 -
Epoch 15/20
60000/60000 [=====] - 8s 137us/step - loss: 0.2183 -
Epoch 16/20
60000/60000 [=====] - 8s 138us/step - loss: 0.2068 -
Epoch 17/20
60000/60000 [=====] - 8s 139us/step - loss: 0.2011 -
Epoch 18/20
60000/60000 [=====] - 8s 137us/step - loss: 0.1886 -
Epoch 19/20
60000/60000 [=====] - 8s 138us/step - loss: 0.1821 -
Epoch 20/20
60000/60000 [=====] - 8s 139us/step - loss: 0.1739 -
```

```
1 score = model_drop.evaluate(X_test, Y_test, verbose=0)
2 print('Test score:', score[0])
3 print('Test accuracy:', score[1])
4
5 fig,ax = plt.subplots(1,1)
6 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
7
8 # list of epoch numbers
9 x = list(range(1,nb_epoch+1))
10
```


```

11 # print(history.history.keys())
12 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
13 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
14
15 # we will get val_loss and val_acc only when you pass the paramter validation_d
16 # val_loss : validation loss
17 # val_acc : validation accuracy
18
19 # loss : training loss
20 # acc : train accuracy
21 # for each key in histry.history we will have a list of length equal to numb
22
23 vy = history.history['val_loss']
24 ty = history.history['loss']
25 plt_dynamic(x, vy, ty, ax)

```



Test score: 0.1093290721397847
Test accuracy: 0.9679



```


1 w_after = model_drop.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)

```

```

11 plt.title("Trained model weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```



Hyper-parameter tuning of Keras models using Sklearn

```

1 from keras.optimizers import Adam,RMSprop,SGD
2 def best_hyperparameters(activ):
3
4     model = Sequential()
5     model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_ini
6     model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean:
7     model.add(Dense(output_dim, activation='softmax'))
8
9
10    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimi

```

```

11
12     return model

1 # https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-1
2
3 activ = ['sigmoid','relu']
4
5 from keras.wrappers.scikit_learn import KerasClassifier
6 from sklearn.model_selection import GridSearchCV
7
8 model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_s
9 param_grid = dict(activ=activ)
10
11 # if you are using CPU
12 # grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
13 # if you are using GPU dont use the n_jobs parameter
14
15 grid = GridSearchCV(estimator=model, param_grid=param_grid)
16 grid_result = grid.fit(X_train, Y_train)

1 print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_)
2 means = grid_result.cv_results_['mean_test_score']
3 stds = grid_result.cv_results_['std_test_score']
4 params = grid_result.cv_results_['params']
5 for mean, stdev, param in zip(means, stds, params):
6     print("%f (%f) with: %r" % (mean, stdev, param))

👤 Best: 0.975633 using {'activ': 'relu'}
0.974650 (0.001138) with: {'activ': 'sigmoid'}
0.975633 (0.002812) with: {'activ': 'relu'}

```

▼ =====Assignment=====

▼ 1. Number of hidden layers = 2 (550,450) + adam optimizer + BN + D

```

1 # for relu layers, directly using he_normal() initializer
2 model_one = Sequential()
3
4 model_one.add(Dense(550, activation='relu', input_shape=(input_dim,), kernel_in
5 model_one.add(BatchNormalization())
6 model_one.add(Dropout(0.6))
7
8 model_one.add(Dense(450, activation='relu', kernel_initializer=he_normal(seed=N
9 model_one.add(BatchNormalization())
10 model_one.add(Dropout(0.5))
11
12 model_one.add(Dense(output_dim, activation='softmax'))
13
14
15 model_one.summary()

```

Model: "sequential_15"

Layer (type)	Output Shape	Param #
dense_40 (Dense)	(None, 550)	431750
batch_normalization_25 (Batch Normalization)	(None, 550)	2200
dropout_25 (Dropout)	(None, 550)	0
dense_41 (Dense)	(None, 450)	247950
batch_normalization_26 (Batch Normalization)	(None, 450)	1800
dropout_26 (Dropout)	(None, 450)	0
dense_42 (Dense)	(None, 10)	4510
Total params: 688,210		
Trainable params: 686,210		
Non-trainable params: 2,000		

```
1 model_one.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
2
3 history = model_one.fit(X_train, Y_train, batch_size=batch_size, epochs=50, verbose=1)
```


Train on 60000 samples, validate on 10000 samples

```
Epoch 1/50
60000/60000 [=====] - 5s 78us/step - loss: 0.4183 - a
Epoch 2/50
60000/60000 [=====] - 3s 54us/step - loss: 0.2105 - a
Epoch 3/50
60000/60000 [=====] - 3s 54us/step - loss: 0.1653 - a
Epoch 4/50
60000/60000 [=====] - 3s 55us/step - loss: 0.1413 - a
Epoch 5/50
60000/60000 [=====] - 3s 55us/step - loss: 0.1264 - a
Epoch 6/50
60000/60000 [=====] - 3s 57us/step - loss: 0.1186 - a
Epoch 7/50
60000/60000 [=====] - 3s 55us/step - loss: 0.1070 - a
Epoch 8/50
60000/60000 [=====] - 3s 54us/step - loss: 0.1018 - a
Epoch 9/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0917 - a
Epoch 10/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0872 - a
Epoch 11/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0873 - a
Epoch 12/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0828 - a
Epoch 13/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0789 - a
Epoch 14/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0766 - a
Epoch 15/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0720 - a
Epoch 16/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0697 - a
Epoch 17/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0677 - a
Epoch 18/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0675 - a
Epoch 19/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0625 - a
Epoch 20/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0568 - a
Epoch 21/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0605 - a
Epoch 22/50
60000/60000 [=====] - 3s 57us/step - loss: 0.0584 - a
Epoch 23/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0555 - a
Epoch 24/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0539 - a
Epoch 25/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0520 - a
Epoch 26/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0530 - a
Epoch 27/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0482 - a
Epoch 28/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0493 - a
Epoch 29/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0470 - a
Epoch 30/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0498 - a
```

```

Epoch 31/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0440 - a
Epoch 32/50
60000/60000 [=====] - 3s 53us/step - loss: 0.0467 - a
Epoch 33/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0425 - a
Epoch 34/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0430 - a
Epoch 35/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0443 - a
Epoch 36/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0413 - a
Epoch 37/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0396 - a
Epoch 38/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0416 - a
Epoch 39/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0376 - a
Epoch 40/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0385 - a
Epoch 41/50
60000/60000 [=====] - 3s 57us/step - loss: 0.0378 - a
Epoch 42/50
60000/60000 [=====] - 3s 57us/step - loss: 0.0365 - a
Epoch 43/50
60000/60000 [=====] - 3s 58us/step - loss: 0.0351 - a
Epoch 44/50
60000/60000 [=====] - 3s 56us/step - loss: 0.0369 - a
Epoch 45/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0351 - a
Epoch 46/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0348 - a
Epoch 47/50
60000/60000 [=====] - 3s 53us/step - loss: 0.0349 - a
Epoch 48/50
60000/60000 [=====] - 3s 54us/step - loss: 0.0363 - a
Epoch 49/50
60000/60000 [=====] - 3s 55us/step - loss: 0.0321 - a
Epoch 50/50
60000/60000 [=====] - 3s 53us/step - loss: 0.0328 - a

```

```

1 %matplotlib inline
2 score = model_one.evaluate(X_test, Y_test, verbose=0)
3 print('Test score:', score[0])
4 print('Test accuracy:', score[1])
5
6 fig,ax = plt.subplots(1,1)
7 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
8
9 # list of epoch numbers
10 x = list(range(1,50+1))
11
12 # print(history.history.keys())
13 # dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
14 # history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_e
15
16 # we will get val_loss and val_acc only when you pass the paramter validation_d
17 # val_loss : validation loss
18 # val_acc : validation accuracy


```

```

19
20 # loss : training loss
21 # acc : train accuracy
22 # for each key in history.history we will have a list of length equal to number of epochs
23
24 vy = history.history['val_loss']
25 ty = history.history['loss']
26 plt_dynamic(x, vy, ty, ax)
27 plt.show()

```

Test score: 0.051541719751636995
Test accuracy: 0.9857

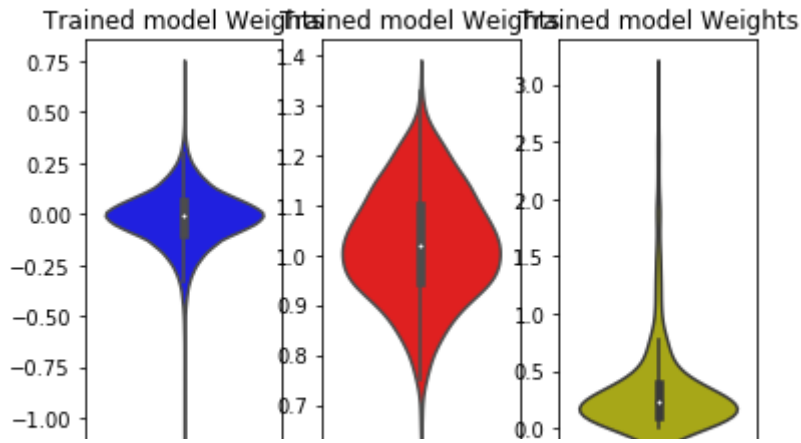


```

1 w_after = model_one.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```





▼ 2. Number of hidden layers = 3 (550,450,350) + adam optimizer + BN

```

1 # for relu layers, directly using he_normal() initializer
2 model_two = Sequential()
3
4 model_two.add(Dense(550, activation='relu', input_shape=(input_dim,), kernel_in
5 model_two.add(BatchNormalization())
6 model_two.add(Dropout(0.6))
7
8 model_two.add(Dense(450, activation='relu', kernel_initializer=he_normal(seed=N
9 model_two.add(BatchNormalization())
10 model_two.add(Dropout(0.5))
11
12 model_two.add(Dense(350, activation='relu', kernel_initializer=he_normal(seed=N
13 model_two.add(BatchNormalization())
14 model_two.add(Dropout(0.5))
15
16 model_two.add(Dense(output_dim, activation='softmax'))
17
18
19 model_two.summary()

```



WARNING:tensorflow:Large dropout rate: 0.6 (>0.5). In TensorFlow 2.x, dropout(
Model: "sequential_13"

Layer (type)	Output Shape	Param #
dense_30 (Dense)	(None, 550)	431750
batch_normalization_17 (Batch Normalization)	(None, 550)	2200
dropout_17 (Dropout)	(None, 550)	0
dense_31 (Dense)	(None, 450)	247950

```
1 model_two.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
2
3 history = model_two.fit(X_train, Y_train, batch_size=batch_size, epochs=25, verbose=1)
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 5s 87us/step - loss: 0.5450 - a

Epoch 2/25

60000/60000 [=====] - 4s 68us/step - loss: 0.2530 - a

Epoch 3/25

60000/60000 [=====] - 4s 67us/step - loss: 0.1952 - a

Epoch 4/25

60000/60000 [=====] - 4s 66us/step - loss: 0.1699 - a

Epoch 5/25

60000/60000 [=====] - 4s 66us/step - loss: 0.1508 - a

Epoch 6/25

60000/60000 [=====] - 4s 70us/step - loss: 0.1379 - a

Epoch 7/25

60000/60000 [=====] - 4s 67us/step - loss: 0.1237 - a

Epoch 8/25

60000/60000 [=====] - 4s 69us/step - loss: 0.1146 - a

Epoch 9/25

60000/60000 [=====] - 4s 67us/step - loss: 0.1133 - a

Epoch 10/25

60000/60000 [=====] - 4s 68us/step - loss: 0.1080 - a

Epoch 11/25

60000/60000 [=====] - 4s 68us/step - loss: 0.1001 - a

Epoch 12/25

60000/60000 [=====] - 4s 68us/step - loss: 0.0951 - a

Epoch 13/25

60000/60000 [=====] - 4s 66us/step - loss: 0.0912 - a

Epoch 14/25

60000/60000 [=====] - 4s 67us/step - loss: 0.0881 - a

Epoch 15/25

60000/60000 [=====] - 4s 68us/step - loss: 0.0847 - a

Epoch 16/25

60000/60000 [=====] - 4s 67us/step - loss: 0.0848 - a


Epoch 17/25

```

1 %matplotlib inline
2 score = model_two.evaluate(X_test, Y_test, verbose=0)
3 print('Test score:', score[0])
4 print('Test accuracy:', score[1])
5
6 fig,ax = plt.subplots(1,1)
7 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
8
9 # list of epoch numbers
10 x = list(range(1,25+1))
11
12 # we will get val_loss and val_acc only when you pass the paramter validation_d
13 # val_loss : validation loss
14 # val_acc : validation accuracy
15
16 # loss : training loss
17 # acc : train accuracy
18 # for each key in histrory.histrory we will have a list of length equal to numb
19
20 vy = history.history['val_loss']
21 ty = history.history['loss']
22 plt_dynamic(x, vy, ty, ax)
23 plt.show()

```

Test score: 0.05603201202231867
 Test accuracy: 0.9837



```

1 w_after = model_two.get_weights()
2
3 h1_w = w_after[0].flatten().reshape(-1,1)
4 h2_w = w_after[2].flatten().reshape(-1,1)
5 out_w = w_after[4].flatten().reshape(-1,1)
6
7
8 fig = plt.figure()
9 plt.title("Weight matrices after model trained")
10 plt.subplot(1, 3, 1)
11 plt.title("Trained model Weights")
12 ax = sns.violinplot(y=h1_w,color='b')
13 plt.xlabel('Hidden Layer 1')
14
15 plt.subplot(1, 3, 2)
16 plt.title("Trained model Weights")
17 ax = sns.violinplot(y=h2_w, color='r')
18 plt.xlabel('Hidden Layer 2 ')
19
20 plt.subplot(1, 3, 3)
21 plt.title("Trained model Weights")
22 ax = sns.violinplot(y=out_w,color='y')
23 plt.xlabel('Output Layer ')
24 plt.show()

```



Trained model Weights Trained model Weights Trained model Weights

3. Number of hidden layers = 5 (650,550,450,350,250) + adam optim Dropout(0.6,0.5,0.5,0.5,0.5)

```

1 # for relu layers, directly using he_normal() initializer
2 model_three = Sequential()
3
4 model_three.add(Dense(650, activation='relu', input_shape=(input_dim,), kernel_
5 model_three.add(BatchNormalization())
6 model_three.add(Dropout(0.6))
7
8 model_three.add(Dense(550, activation='relu', kernel_initializer=he_normal(seed:
9 model_three.add(BatchNormalization())
10 model_three.add(Dropout(0.5))
11
12 model_three.add(Dense(450, activation='relu', kernel_initializer=he_normal(seed:
13 model_three.add(BatchNormalization())
14 model_three.add(Dropout(0.5))
15
16 model_three.add(Dense(350, activation='relu', kernel_initializer=he_normal(seed:
17 model_three.add(BatchNormalization())
18 model_three.add(Dropout(0.5))
19
20 model_three.add(Dense(250, activation='relu', kernel_initializer=he_normal(seed:
21 model_three.add(BatchNormalization())
22 model_three.add(Dropout(0.5))
23
24 model_three.add(Dense(output_dim, activation='softmax'))
25
26
27 model_three.summary()
```



Model: "sequential_14"

Layer (type)	Output Shape	Param #
dense_34 (Dense)	(None, 650)	510250
batch_normalization_20 (Batch Normalization)	(None, 650)	2600
dropout_20 (Dropout)	(None, 650)	0
dense_35 (Dense)	(None, 550)	358050
batch_normalization_21 (Batch Normalization)	(None, 550)	2200
dropout_21 (Dropout)	(None, 550)	0
dense_36 (Dense)	(None, 450)	247950
batch_normalization_22 (Batch Normalization)	(None, 450)	1800
dropout_22 (Dropout)	(None, 450)	0
dense_37 (Dense)	(None, 350)	157850

```
1 model_three.compile(optimizer='adam', loss='categorical_crossentropy', metrics=
2
3 history = model_three.fit(X_train, Y_train, batch_size=batch_size, epochs=25, v
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/25

60000/60000 [=====] - 7s 124us/step - loss: 0.8781 - a

Epoch 2/25

60000/60000 [=====] - 6s 92us/step - loss: 0.3221 - a

Epoch 3/25

60000/60000 [=====] - 6s 92us/step - loss: 0.2472 - a

Epoch 4/25

60000/60000 [=====] - 5s 91us/step - loss: 0.2086 - a

Epoch 5/25

60000/60000 [=====] - 6s 93us/step - loss: 0.1839 - a

Epoch 6/25

60000/60000 [=====] - 6s 93us/step - loss: 0.1685 - a

Epoch 7/25

60000/60000 [=====] - 5s 92us/step - loss: 0.1554 - a

Epoch 8/25

60000/60000 [=====] - 6s 92us/step - loss: 0.1438 - a

Epoch 9/25

60000/60000 [=====] - 6s 92us/step - loss: 0.1350 - a

Epoch 10/25

60000/60000 [=====] - 5s 91us/step - loss: 0.1313 - a

Epoch 11/25

60000/60000 [=====] - 6s 93us/step - loss: 0.1240 - a

Epoch 12/25

60000/60000 [=====] - 6s 92us/step - loss: 0.1179 - a

Epoch 13/25

60000/60000 [=====] - 6s 92us/step - loss: 0.1129 - a

Epoch 14/25

60000/60000 [=====] - 6s 94us/step - loss: 0.1075 - a

Epoch 15/25

60000/60000 [=====] - 6s 92us/step - loss: 0.1045 - a

Epoch 16/25

60000/60000 [=====] - 6s 93us/step - loss: 0.1004 - a

Epoch 17/25

```

1 %matplotlib inline
2 score = model_three.evaluate(X_test, Y_test, verbose=0)
3 print('Test score:', score[0])
4 print('Test accuracy:', score[1])
5
6 fig,ax = plt.subplots(1,1)
7 ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
8
9 # list of epoch numbers
10 x = list(range(1,25+1))
11
12 # we will get val_loss and val_acc only when you pass the paramter validation_d
13 # val_loss : validation loss
14 # val_acc : validation accuracy
15
16 # loss : training loss
17 # acc : train accuracy
18 # for each key in histrory.histrory we will have a list of length equal to numb
19
20 vy = history.history['val_loss']
21 ty = history.history['loss']
22 plt_dynamic(x, vy, ty, ax)
23 plt.show()

```